CLIENT PACK · DOCUMENT 06

Every website section

For each menu: what it is used for, the inputs you fill, how to use it, and how it connects.

Connection flow and every website section

Do the steps in this order. Logs stay empty until a provider and a virtual key exist. Sample values are for training.

Order	Do this first
1	Settings → Security: save SMTP and confirm the test mail.
2	Settings → Client Settings. Leave Proxy off unless the server needs an office proxy.
3	Models → Model Providers, then Model Catalog.
4	Governance → Teams, Users, then Virtual Keys.
5	Routing Rules, Budgets & Limits, Pricing. Then Guardrails for API text.
6	One API call, then Observability → LLM Logs.
7	Browser AI: Target Websites, Guard Rules, Setup, then Guard Agents and Prompt Logs.
8	MCP Catalog, MCP Auth Config, and Tool Groups only if a tool server is required.
9	Prompt Repository, then Prompt Deployments. Skills Repository when you publish a skill.
10	Alert Channels after a budget or circuit breaker exists. Connectors when logs must leave the server.

An API call connects like this: virtual key, then budget, then routing, then guardrails, then the model provider, then LLM Logs. A browser prompt connects like this: target website, then the laptop agent, then the guard rule, then Prompt Logs. That path does not use a virtual key.

1. Observability

Dashboard

Open: /workspace/dashboard

What this is used for

Shows company totals for API traffic: requests, tokens, cost, and latency.

Inputs you fill

Input	What you enter
Date range	Last 7 days
No create fields	This page is read-only

How to use it

Open it after at least one virtual-key call has succeeded. Read the cards. A high cost sends you to LLM Logs for that model or key.

How it connects

Reads LLM Logs. Cost uses Pricing Overrides and Model Settings. Empty means Model Providers and Virtual Keys are not in use yet.

LLM Logs

Open: /workspace/logs

What this is used for

One row for every API call made with a virtual key.

Inputs you fill

Input	What you enter
Date	Today
Virtual key	Support bot — production
Model	gpt-4.1-mini
Status	success or error

How to use it

Filter, open a row, and read the prompt, answer, tokens, cost, and latency. Export if you need a sheet. Clear logs or recalculate cost only when an admin means to.

How it connects

A row is written after Virtual Keys, Routing, Guardrails, and Model Providers finish a call. A 401 points at the provider key. A budget error points at Virtual Keys or Budgets & Limits.

MCP Logs

Open: /workspace/mcp-logs

What this is used for

One row for every tool call made through MCP Gateway.

Inputs you fill

Input	What you enter
Client	docs-server
Tool	search_docs
User	anitha.user
Time	Today

How to use it

Filter and open the row to see the tool arguments and the result.

How it connects

Appears only after MCP Catalog is connected, MCP Auth Config is saved, a Tool Group allows the tool, and the feature flag mcp.tool_execution is on.

Connectors

Open: /workspace/observability

What this is used for

Copies logs and metrics to a system the company already runs.

Inputs you fill

Input	What you enter
Connector	Datadog, BigQuery, Kafka, Pub/Sub, New Relic, Prometheus, or OpenTelemetry
URL or site	The vendor endpoint
API key or dataset	The secret from that vendor account

How to use it

Add connector, paste the vendor values, save, and send one test chat. Confirm the same request in LLM Logs and in the vendor.

How it connects

Source is LLM Logs and metrics. This does not send the short budget alerts. Those go to Alert Channels.

Logs Settings

Open: /workspace/config/logging

What this is used for

Chooses what API log text is stored and for how long.

Inputs you fill

Input	What you enter
Content logging	On for a pilot, off if prompts must not be stored
Retention days	30

How to use it

Set the two values and save. New LLM Log rows follow this. Old rows stay until retention deletes them.

How it connects

Controls LLM Logs and the Observability Dashboard. It does not control Browser AI Prompt Logs.

2. Browser AI

Overview

Open: /workspace/browser-ai?tab=overview

What this is used for

Counts monitored sites, active rules, recent agents, and recent decisions.

Inputs you fill

Input	What you enter
No create fields	Read the counts after the other Browser AI tabs are filled

How to use it

Use it as the check that the rollout is alive. Zero monitored sites means Target Websites is empty. Zero agents means Setup has not reported a laptop.

How it connects

Summarises Target Websites, Guard Rules, Prompt Logs, and Guard Agents.

Target Websites

Open: /workspace/browser-ai?tab=targets

What this is used for

The public AI sites the laptop is allowed to watch.

Inputs you fill

Input	What you enter
Domain	chatgpt.com
Platform name	ChatGPT
Host role	ui, chat, or file
Monitored	Yes
Block entire website	No
Status	MONITORED

How to use it

Add one row or import a CSV with those columns. Add the chat and file hosts when the vendor uses more than one domain. Pause with status PAUSED.

How it connects

Gateway Guard downloads this list. A Guard Rule does nothing on a domain that is not here.

Guard Rules

Open: /workspace/browser-ai?tab=rules

What this is used for

What the laptop does when a prompt or upload on a target site matches.

Inputs you fill

Input	What you enter
Rule name	India PAN
Rule type	regex, or ai_bot typed in the screen
Pattern	[A-Z]{5}[0-9]{4}[A-Z]
Severity	CRITICAL, HIGH, or MEDIUM
Action	BLOCK, WARN, or REDACT
Warning message	Personal ID numbers cannot be sent to a public AI site.
Active	TRUE

How to use it

Create or import. BLOCK stops the send. WARN shows the message and sends the original. REDACT shows the message and sends the match as [REDACTED]. Test an ai_bot rule before Active is TRUE. AI bot rules cannot be imported.

How it connects

Runs on text sent to a Target Website. The result is a Prompt Log. This is not Guardrails. Guardrails check API calls.

Prompt Logs

Open: /workspace/browser-ai?tab=logs

What this is used for

Every browser prompt the agent checked.

Inputs you fill

Input	What you enter
Date	Today
Action	Blocked, Warned, Redacted, Allowed, SiteBlocked, or Bot Answered
Search	The laptop hostname

How to use it

Filter and open the row. You see the prompt, the rule, the action, and the device.

How it connects

Written by Gateway Guard after a Guard Rule runs. Match the hostname on Guard Agents.

Search Logs

Open: /workspace/browser-ai?tab=search-logs

What this is used for

Search queries the agent recorded, including private windows.

Inputs you fill

Input	What you enter
Engine	Google
Browser	Chrome, Edge, or Brave
Private window	Yes or no
Query	The text the employee searched

How to use it

Filter and open a row for the query and any result that was opened.

How it connects

Same laptop agent as Prompt Logs. It is not an LLM Log.

Setup

Open: /workspace/browser-ai?tab=setup

What this is used for

Download the laptop package and publish Guard code updates.

Inputs you fill

Input	What you enter
Package	Gateway_Guard_Setup.exe for Windows, or the Mac package
Rebuild	Use Rebuild and Publish only after Guard code on the server changed

How to use it

Download the package and give it to IT. IT runs it once per laptop and leaves start at login checked.

How it connects

The package installs Gateway Guard, which force-installs YP-Raksha and then appears under Guard Agents.

Guard Agents

Open: /workspace/browser-ai?tab=agents

What this is used for

Every laptop and every office proxy that has reported in.

Inputs you fill

Input	What you enter
Search	DESKTOP-EXAMPLE
Contact email	it@example.com
Uninstall key	A secret kept in the IT vault, not in email

How to use it

Confirm status active and a recent last-seen time. Remote uninstall moves the device to uninstall_pending, then uninstalled.

How it connects

The agent appears only after Setup has run and the laptop can reach this server. Prompt Logs use the same hostname.

Guard Insights

Open: /workspace/browser-ai?tab=telemetry

What this is used for

Totals of allowed, blocked, warned, and redacted events per laptop.

Inputs you fill

Input	What you enter
No create fields	Read the totals, then open Prompt Logs for the loudest hostname

How to use it

Use a high Blocked count as a reason to open that laptop’s Prompt Logs before you edit the rule.

How it connects

Adds up Prompt Logs. It does not change Guard Rules.

3. Models

Model Providers

Open: /workspace/providers

What this is used for

Attaches the company’s account at OpenAI, Anthropic, Azure, Bedrock, Gemini, or another vendor. Admin only for the secret.

Inputs you fill

Input	What you enter
Provider	OpenAI
Key name	openai-prod-1
API key	The vendor key
Azure or Bedrock extras	Endpoint or cloud credentials, only if that form asks

How to use it

Add provider, paste the key, save. You may store more than one key for the same vendor. A keyless provider needs no key.

How it connects

Model Catalog lists models from these keys. Routing sends calls here. The virtual key never contains this vendor key.

Model Catalog

Open: /workspace/model-catalog

What this is used for

The models UnifAI is allowed to offer.

Inputs you fill

Input	What you enter
Model	gpt-4.1-mini
Enabled	Yes for models under contract, no for the rest

How to use it

Search, enable, and disable. Open a model to see its price.

How it connects

Fed by Model Providers. Virtual Keys and Routing Rules should name only enabled models.

Budgets & Limits

Open: /workspace/model-limits

What this is used for

A spend cap and a request rate on one model, separate from a virtual-key budget.

Inputs you fill

Input	What you enter
Model	gpt-4.1-mini
Budget	500
Rate	60 requests per minute

How to use it

Open the model, set the ceiling, save. The first ceiling that is hit stops the call.

How it connects

Checked with the Virtual Key, the Team, and the Business Unit. The error is in LLM Logs.

Complexity Router

Open: /workspace/complexity-router

What this is used for

Sends an easy prompt to a cheaper model and a hard prompt to a stronger model.

Inputs you fill

Input	What you enter
Enabled	Yes
Low complexity model	gpt-4.1-mini
High complexity model	gpt-4.1
Feature flag	routing.complexity on

How to use it

Save the map. Send a short hello and confirm LLM Logs shows the small model.

How it connects

Runs before the provider call. Routing Rules can still force one model. Both models must be enabled in the catalog.

Routing Rules

Open: /workspace/routing-rules

What this is used for

Company policy: which provider and model for a team, user, key, or model name.

Inputs you fill

Input	What you enter
Condition	Team is Support
Target model	gpt-4.1-mini
Order	Specific rule above the general rule

How to use it

Add the rule and save. Send a Support-key call and confirm the model in LLM Logs.

How it connects

Reads the virtual key, user, and team. Sends the call to a Model Provider. Use Complexity Router for automatic difficulty. Use this screen for a fixed policy.

Circuit Breaker

Open: /workspace/circuit-breaker

What this is used for

Stops sending traffic to a provider that is failing.

Inputs you fill

Input	What you enter
Failures before open	5 in a minute
Wait before retry	30 seconds
Feature flag	routing.circuit_breaker on

How to use it

Save. You need a second provider key if traffic must continue while one vendor is paused.

How it connects

Watches errors from Model Providers. Adaptive Routing and Routing Rules should use the healthy provider. LLM Logs show the errors that opened it.

Pricing Overrides

Open: /workspace/custom-pricing/overrides

What this is used for

Your contract price, used instead of the default price.

Inputs you fill

Input	What you enter
Model	gpt-4.1-mini
Input price	Your contract price per million tokens
Output price	Your contract price per million tokens

How to use it

Save. After a change, an admin can recalculate cost from LLM Logs.

How it connects

Dashboard, LLM Logs, and budgets use this price when it is set. It wins over Model Settings when both exist.

Model Settings

Open: /workspace/custom-pricing

What this is used for

Default parameters and the base price for a model.

Inputs you fill

Input	What you enter
Model	gpt-4.1-mini
Temperature	0.2 for a stable support answer
Base price	Set here only if you are not using a Pricing Override

How to use it

Open the model, set the defaults, save.

How it connects

Applied on the way to the provider unless the client sends its own value. Pricing Overrides win on price.

4. MCP Gateway

MCP Catalog

Open: /workspace/mcp-registry

What this is used for

The MCP servers UnifAI will dial.

Inputs you fill

Input	What you enter
Name	docs-server
URL or command	https://mcp.example.com/sse

How to use it

Add the client, save, and reconnect if the first dial fails. Status should become connected.

How it connects

Tool Groups pick tools from here. MCP Auth Config stores the header. MCP Logs show the calls. Flag mcp.tool_execution must be on.

MCP Library

Open: /workspace/mcp-registry/library

What this is used for

A ready list of known MCP servers. It is a template, not a live connection.

Inputs you fill

Input	What you enter
Server	The documented server you need
Sync	Refresh the list if it looks stale

How to use it

Add the chosen server into MCP Catalog, then fill the real URL and credentials there.

How it connects

Library does not dial anything. Catalog does.

Tool Groups

Open: /workspace/mcp-tool-groups

What this is used for

A named set of tools one virtual key may call.

Inputs you fill

Input	What you enter
Group name	Support tools
Tools	search_docs only
Attached to	The Support virtual key

How to use it

Create the group, tick only those tools, and attach it to the key.

How it connects

A key without the group cannot call the tool. Calls appear in MCP Logs.

Auth Sessions

Open: /workspace/mcp-sessions

What this is used for

Live sign-ins between a person and an MCP server.

Inputs you fill

Input	What you enter
User	anitha.user
Server	docs-server
State	success or the failure reason

How to use it

Open a session to see who connected. Disconnect a session that should stop.

How it connects

Created after MCP Auth Config and, when required, OAuth Grants. Feature flag mcp.oauth_grants turns on the consent screen.

OAuth Grants

Open: /workspace/oauth-grants

What this is used for

The consent a person gave an MCP server to use an external account.

Inputs you fill

Input	What you enter
User	anitha.user
Server	docs-server
Action	Revoke on the day the person leaves

How to use it

Review the list. Revoke instead of leaving a grant in place.

How it connects

Sits between the user and the external service. Auth Sessions show whether the sign-in worked.

MCP Settings

Open: /workspace/mcp-settings

What this is used for

Gateway-wide timeouts and switches for every MCP server.

Inputs you fill

Input	What you enter
Timeout	Raise it only when a known server is slow

How to use it

Change the value and save. Then reconnect the client from MCP Catalog.

How it connects

Applies to every Catalog server. The URL of each server stays on the catalog entry.

Plugins

Open: /workspace/plugins

What this is used for

Turns the processing steps on the API path on or off, and sets their order.

Inputs you fill

Input	What you enter
Keep on	governance and logging
Leave off	Any plugin you do not recognise

How to use it

Enable, disable, and reorder only when you know which step must run first. Send one test chat and confirm LLM Logs still has a row.

How it connects

Governance off stops budget checks. Guardrails off stops API guardrail rules. This page is not the laptop Guard.

5. Governance

Virtual Keys

Open: /workspace/governance/virtual-keys

What this is used for

The secret an application sends. It is not a website login and it is not the vendor key.

Inputs you fill

Input	What you enter
Name	Support bot — production
Team or user	Support
Budget	200
Rate	120 per minute
Models	gpt-4.1-mini

How to use it

Create, copy the key once into the application secret store, and rotate it when staff leave.

How it connects

The app calls /v1/chat/completions with Authorization: Bearer and this key. Budgets, Routing, Guardrails, and Model Providers run on that call. LLM Logs show it.

Users

Open: /workspace/governance/users

What this is used for

Website accounts and who may open which menu.

Inputs you fill

Input	What you enter
Username	anitha.user
Email	anitha.user@example.com
Password	A password with 8 characters, upper, lower, number, and a symbol
Role	user, sub_admin, admin, or a custom role
Allowed sections	prompt-repository, and browser-ai only for security staff
Status	Approve a pending sign-up, or disable a leaver

How to use it

Create the user or approve the pending row. The person then signs in at /login.

How it connects

Allowed sections decide the menu. A virtual key can be assigned to this user. SMTP must already work if they signed up with an email code.

Teams

Open: /workspace/governance/teams

What this is used for

A group of users with one shared budget.

Inputs you fill

Input	What you enter
Name	Support
Members	anitha.user
Budget	200

Input	What you enter
Rate	120 per minute

How to use it

Create the team, add the people, set the ceiling, save.

How it connects

Put the team on a Virtual Key and on a Routing Rule. A Business Unit above the team can stop spend even when the team budget remains.

Business Units

Open: /workspace/governance/business-units

What this is used for

A department budget over several teams.

Inputs you fill

Input	What you enter
Name	Customer Operations
Teams	Support
Budget	2000
Rate	The department ceiling

How to use it

Create the unit, attach the teams, save.

How it connects

Sits above Teams. The unit ceiling stops every team under it.

Customers

Open: /workspace/governance/customers

What this is used for

An external client you report and bill separately.

Inputs you fill

Input	What you enter
Name	Northwind
Teams	Support

How to use it

Create the customer and link the teams that do that client’s work.

How it connects

Reporting and budgets. The application still uses a Virtual Key.

User Provisioning

Open: /workspace/scim

What this is used for

Connects Azure AD or Okta so joiners and leavers update Users automatically.

Inputs you fill

Input	What you enter
SCIM base URL	https://unifai.example.com/api/scim/v2
Token	The token this page shows
Map	User name and email

How to use it

Paste the URL and token into the identity provider. A person set inactive there becomes disabled here.

How it connects

Writes the same Users records. It does not install Gateway Guard and it does not replace SMTP.

Roles & Permissions

Open: /workspace/governance/rbac

What this is used for

What a role may do on the API.

Inputs you fill

Input	What you enter
Role name	Auditor
Data scope	all-data or own-data
Allow	Logs Read and View, Audit Logs Read and View, Dashboard Read and View
Do not allow	Delete, provider keys, Settings

How to use it

Create the custom role and assign it on the user. Leave admin, sub_admin, and user in place unless you must narrow them.

How it connects

Every API call checks this, then the allowed sections on the user. Browser AI is under the Logs resource. The Users page itself stays with admin and sub_admin.

Access Profiles

Open: /workspace/governance/access-profiles

What this is used for

A reusable bundle of budget, model, and key limits for one job.

Inputs you fill

Input	What you enter
Name	Contractor
Budget	25
Models	gpt-4.1-mini only
State	Active

How to use it

Create, activate, and apply it to those users. Clone it for a similar job.

How it connects

Applying it updates those users’ virtual-key limits.

Audit Logs

Open: /workspace/audit-logs

What this is used for

Who changed configuration. You do not type these rows.

Inputs you fill

Input	What you enter
User	priya.admin
Date	Today
Action	The create or update you just saved

How to use it

Filter after an incident. Do not delete the rows as a habit.

How it connects

Records changes on Users, Virtual Keys, Guard Rules, Model Providers, and Settings.

6. Guardrails

Rules

Open: /workspace/guardrails/configuration

What this is used for

Checks prompts and answers on API calls that use a virtual key.

Inputs you fill

Input	What you enter
Detection	PII and secrets, prompt injection, toxicity, or grounding
Action	Block or redact the API call
Feature flags	guardrails.pii_masking and the other guardrails flags you rely on

How to use it

Create the rule, point it at a Rule Provider if it needs one, and save. Send a test API call that contains a card number and read LLM Logs.

How it connects

Runs on the API path. It does not write a Browser AI Prompt Log. Browser checks stay in Guard Rules.

Rule Providers

Open: /workspace/guardrails/providers

What this is used for

The safety engine a guardrail rule calls. This is not the chat vendor.

Inputs you fill

Input	What you enter
Provider	The safety product you licensed
Key or endpoint	That product’s credential

How to use it

Add it, save, then select it on the rule.

How it connects

Model Providers call the chat model. This page calls the checker.

Cluster Config

Open: /workspace/cluster

What this is used for

More than one UnifAI server sharing configuration.

Inputs you fill

Input	What you enter
Peers	Leave empty on a single server
Cluster secret	The same secret on every node, only when you have a pair

How to use it

On one server, do not add peers. On a pair, each node lists the other.

How it connects

A virtual key created on one node must exist on the others. This does not replace the database.

7. Adaptive Routing

Dashboard

Open: /workspace/adaptive-routing

What this is used for

How API traffic is split across provider keys right now. This is not the Observability Dashboard.

Inputs you fill

Input	What you enter
No create fields	Read the weight and the health of each route

How to use it

If one provider takes every call, open Adaptive Routing Settings or Routing Rules.

How it connects

Uses the Model Provider pool. Circuit Breaker removes a failing provider from this split.

Settings

Open: /workspace/adaptive-routing/settings

What this is used for

The load-balancing policy when more than one key can serve the model.

Inputs you fill

Input	What you enter
Strategy	The split you want
Weight openai-prod-1	70
Weight openai-prod-2	30
Fallback flag	routing.fallback_cascade on if a 429 or 500 must jump to the next provider

How to use it

Set the weights and save.

How it connects

Routing Rules win when a rule names one provider. Circuit Breaker removes a sick key from the weights.

8. Prompt Repository and Skills Repository

Prompt Repository

Open: /workspace/prompt-repo

What this is used for

Saved prompt text for people and for applications.

Inputs you fill

Input	What you enter
Folder	Support
Name	Customer reply — refund tone
Text	Reply in a calm tone. Do not ask for card numbers or government IDs.
Tags	support, refund

How to use it

Create the folder, add the prompt, save. Limit folders on the user if another team must not see it.

How it connects

Does not choose a model. Prompt Deployments publishes a version. The call still needs a Virtual Key.

Skills Repository

Open: /workspace/skills-repo

What this is used for

Shared skill instructions the company published.

Inputs you fill

Input	What you enter
Name	Refund tone
Body	Three short sentences. No legal promises. Do not ask for a PAN or a card number.

How to use it

Add the skill and save. People with this menu can browse it.

How it connects

Served to clients that call the skills API. It is not a Guard Rule and not a virtual key.

9. Settings

Client Settings

Open: /workspace/config/client-settings

What this is used for

Limits for incoming API calls.

Inputs you fill

Input	What you enter
Allowed origins	https://unifai.example.com
Max body size	100 MB unless you must raise it
Require auth	On

How to use it

Save. Applications must call from an allowed origin and must still send a valid virtual key.

How it connects

This is not Target Websites. AI sites the laptop watches are under Browser AI.

Compatibility

Open: /workspace/config/compatibility

What this is used for

Converts a client’s parameters into the shape a provider expects.

Inputs you fill

Input	What you enter
Conversion	On when an OpenAI-style client calls a non-OpenAI model through the gateway

How to use it

Turn on only the conversion you need and save.

How it connects

Sits between the virtual-key request and Model Providers. Routing still chooses the provider.

Caching

Open: /workspace/config/caching

What this is used for

Reuses an earlier answer for a repeated request.

Inputs you fill

Input	What you enter
Direct cache	On for a FAQ bot
Semantic cache	On only if similar questions should share an answer
TTL	10 minutes
Feature flag	cache.responses on, or every call still goes to the provider

How to use it

Save. Turn the flag off while you are debugging a prompt.

How it connects

A cache hit does not call Model Providers. LLM Logs should show the hit.

Security

Open: /workspace/config/security

What this is used for

Sign-in protection, and the SMTP mail server for codes and notices.

Inputs you fill

Input	What you enter
Lockout	On
SMTP enabled	Yes
Host	smtp.example.com
Port	587
Username	mailer@example.com
From email	noreply@example.com
From name	UnifAI
Use TLS	Yes
Notify on user create	Yes

How to use it

Save SMTP, then send a test to priya.admin@example.com. Sign-up codes and password reset start only after that test arrives.

How it connects

Used by Sign Up, Forgot password, and Forgot username. An email Alert Channel also needs this. Virtual keys do not use this login.

API Keys

Open: /workspace/config/api-keys

What this is used for

The same vendor secrets as Model Providers, opened from Settings.

Inputs you fill

Input	What you enter
Key name	openai-prod-1
New secret	The replacement vendor key

How to use it

Rotate here when you are already in Settings. Prefer Model Providers for the full provider form. Send one test chat and read LLM Logs.

How it connects

The next routed call uses the new secret.

Performance Tuning

Open: /workspace/config/performance-tuning

What this is used for

Server buffers and pool sizes.

Inputs you fill

Input	What you enter
Leave defaults	On a new install
Change	Only after latency on the Observability Dashboard is high while the vendor itself is fast

How to use it

Save and watch the Dashboard.

How it connects

Affects how fast calls are handled. It does not choose the model.

Feature Flags

Open: /workspace/config/feature-flags

What this is used for

On and off switches for whole features.

Inputs you fill

Input	What you enter
browser.prompt_guard	On, or laptop prompts are not stored
browser.file_scan	On if uploads must be checked
cache.responses	Off until answers are trusted
routing.complexity	On only if Complexity Router is in use
routing.circuit_breaker	On only if Circuit Breaker is in use
routing.fallback_cascade	On if a 429 or 500 must jump provider
mcp.tool_execution	On only if MCP tools must run
mcp.oauth_grants	On if MCP must show a consent screen

How to use it

Add the preset and set it on or off. A flag gates the screen it names.

How it connects

browser.prompt_guard off means Prompt Logs stay empty. routing.circuit_breaker off means the Circuit Breaker page does not enforce.

10. Beside the menu, in the same order

Proxy

Open: /workspace/config/proxy

What this is used for

The office HTTP proxy the UnifAI server uses for outbound calls. Not Gateway Guard and not the laptop PAC.

Inputs you fill

Input	What you enter
Enabled	No, unless the server cannot reach vendors directly
URL	http://proxy.example.com:8080
Enable for inference	Yes, only when the proxy is on
Enable for SCIM	Yes only if the identity provider is reached through that proxy
No proxy	localhost,127.0.0.1

How to use it

Leave it off on a normal pilot. If you turn it on, save and confirm one chat still succeeds in LLM Logs.

How it connects

Model Providers use it for vendor calls when inference is enabled. User Provisioning uses it when SCIM is enabled.

Alert Channels

Open: /workspace/alert-channels

What this is used for

A short notice when a circuit breaker trips or a budget or rate limit is exhausted. The same incident is sent at most once every 10 minutes.

Inputs you fill

Input	What you enter
Name	Ops Slack
Type	Webhook, Slack, Email, or PagerDuty
URL, email, or routing key	The Slack webhook URL, or the address, or the PagerDuty key
Enabled	Yes

How to use it

Add the channel after a budget or a circuit breaker exists. Press test. Email type needs SMTP already tested.

How it connects

Connectors copy full logs. This page sends the short incident only.

MCP Auth Config

Open: /workspace/mcp-auth-config

What this is used for

The header and per-user credential for an MCP server that is already in the catalog.

Inputs you fill

Input	What you enter
Client	docs-server
Header name	Authorization
Header value	The server token
Per-user header keys	Authorization, when each person has their own credential

How to use it

If the page says there are no clients, add the server in MCP Catalog first. Save headers, then save the per-user list. Reconnect the catalog client if the state stays failed.

How it connects

Auth Sessions and OAuth Grants show the result. MCP Logs show the tool call.

Prompt Deployments

Open: /workspace/prompt-repo

What this is used for

Publishes one prompt version to production, staging, or preview.

Inputs you fill

Input	What you enter
Prompt	Customer reply — refund tone
Version	1, created in Prompt Repository first
Environment	production, staging, or preview
Enabled	Yes

How to use it

Select the prompt, deploy the version, and leave the switch on. Turn it off to stop that environment without deleting the prompt.

How it connects

The application sends header x-uf-prompt-environment: production, or x-uf-dim-environment, plus the virtual key. The model call still follows Routing and Model Providers.