

UnifAI - Technical Documentation

Audience: developers, architects, integrators

Scope: architecture, modules, data model, HTTP surface, plugins, Browser AI internals, UI, security model

Baseline: repo commit **16d4c3e**

1. System overview

[illegible]

Two data planes exist in the same binary:

- 1. Data plane** (/v1/*, pass-through integrations) - request/response proxying to AI providers,

with governance, guardrails, caching, logging, tracing.

2. **Control plane (/api/*)** - the dashboard: configuration, RBAC, observability and the whole

Browser AI guardrail control plane.

2. Technology stack

Layer	Technology	Source of truth
Gateway	Go 1.26.4, valyala/fasthttp , fasthttp/router , automaxprocs	go.work , transports/unifai-http/main.go
Persistence	GORM; Postgres / MySQL / SQLite (CGO go-sqlite3 , sqlite_static)	framework/configstore , framework/logstore
Provider SDKs	Vendor HTTP clients + official SDKs	core/providers/*
Plugins	Go interfaces, optional .so dynamic loading (SharedObjectPluginLoader)	plugins/ , server/plugins.go
MCP	Model Context Protocol client/server, OAuth2 sessions	core/mcp , handlers/mcp*.go
Frontend	React 19, TypeScript 5.9, Vite 8, TanStack Router/Table, Redux Toolkit, Radix UI, Tailwind 4, Monaco, Recharts, Zod	ui/package.json
Guard agent	Python 3.11+, mitmproxy addon, PyInstaller, Inno Setup (Windows), pkg/dmg scripting (macOS)	apps/browser-guard/**

Layer	Technology	Source of truth
Deploy	Docker (multi-stage), docker compose, Helm (Postgres, Redis, Qdrant, Weaviate, HPA, Ingress)	deploy/**, docker-compose.yml
Tooling	Make targets, gotestsum, Playwright, newman, oxlint/oxfmt	Makefile

3. Repository layout

```
core/ Go module: the AI engine providers/<vendor>/ one package per provider (chat, responses, images, audio, video, batch...) schemas/ shared types (requests, responses, plugin interfaces, errors) mcp/ MCP client, code mode, credential store keyselectors/ weighted/random key selection strategies network/ HTTP client utilities (retries, timeouts, proxies) framework/ Shared infrastructure configstore/ ConfigStore + WorkspaceStore; tables/ has every GORM model logstore/ LLM/MCP/Browser AI log storage, aggregations, matviews, cleaners rbac/ Resource×Operation enforcement, route → requirement mapping encrypt/ AES encryption at rest for secrets in DB cluster/, loadbalancer/, routing/, circuitbreaker/, kvstore/, featureflags/ connectors/, objectstore/, vectorstore/, oauth2/, mailer/, alerts/, tracing/ migrator/, mysqlconn/, postgresconn/ DB drivers + schema migration helpers plugins/ telemetry, prompts, logging, governance, otel, semanticcache, compat, maxim, guardrails, connectors, modelcatalogresolver, mocker, jsonparser transports/unifai-http/ The server binary main.go flags, logger, version banner, bootstrap/start server/server.go UnifAIHTTPServer: bootstrap, middleware assembly, graceful shutdown server/plugins.go built-in plugin instantiation + ordering handlers/*.go every /api and /v1 handler (one file per feature area) integrations/*.go pass-through provider surfaces websocket/, ui/ WS pool/sessions; embedded UI assets ui/ React dashboard (Vite). pages under ui/app/workspace/** apps/browser-guard/ Python Guard agent + mitmproxy addon + installers + release artifacts configs/ config.json, mcp-library.json, model-parameters.json, pricing.json deploy/ docker/, helm/unifai/ docs/ this documentation set
```

4. Core module (core/)

4.1 Providers

Each directory under **core/providers/<vendor>** implements the shared provider contract from **core/schemas**. A provider package contains files such as **chat.go, responses.go, images.go, speech.go, transcription.go, videos.go, batch.go, models.go, errors.go, types.go, utils.go**. Capability differences that matter (from the files present):

Provider(s)	Notable capability files
openai	chat, responses, responseslifecycle, text, images, speech, transcription, videos, batch, files, embedding, realtime, websocket, large_payload, passthrough_usage
anthropic, gemini	chat, responses, count_tokens / models, embedding (gemini), types
vertex	chat, batch, cachedcontents, count_tokens, embedding, rerank, models
bedrock, bedrockmantle	chat, responses, embedding, images, batch
replicate, runware, runway	chat, images, videos, files
mistral	adds ocr.go on top of chat/transcription/cachedcontents
elevenlabs	audio-only

Provider(s)	Notable capability files
<code>ollama</code> , <code>sgl</code> , <code>vllm</code> , <code>openrouter</code> , <code>parasail</code> , <code>opencode</code> , <code>groq</code> , <code>cerebras</code> , <code>deepseek</code> , <code>fireworks</code> , <code>huggingface</code> , <code>nebius</code> , <code>perplexity</code> , <code>xai</code> , <code>cohere</code>	vendor subsets (mostly OpenAI-compatible shapes)
<code>openaicompat</code>	generic OpenAI-compatible provider + <code>registry.go</code> for many compatible endpoints
<code>azure</code>	Azure OpenAI, deployment-aware

Shared helpers live in `core/providers/utils`: `decompression.go`, `sse.go`, `stream.go`, `streamterminaldetector.go`, `pagination.go`, `images.go`, `audio.go`, `file.go`, `videos.go`, `bodysigner.go` (request body signing for AWS-style auth), `modelparamscache.go`, `largerresponse.go`, `passthrough_stream.go`, `fetch.go`, `models.go`, `utils.go`.

4.2 Schemas

`core/schemas` holds the types shared by the gateway and every plugin: chat / responses / text / embedding / image / audio / video request-response types, `UnifAIConfig`, `UnifAIContext`, `ProviderConfig`, `NetworkConfig`, the plugin interfaces (`BasePlugin`, `LLMPlugin`, `MCPPlugin`, `ObservabilityPlugin`, `HTTPTransportPlugin`, `UnifAIHTTPMiddleware`), error types and log level/output enums.

4.3 MCP client

`core/mcp` implements the Model Context Protocol client used by the MCP Gateway: connection lifecycle, tool discovery, tool execution, per-user credentials (`credstore`), "code mode" (`codemode`) and utilities. The HTTP layer wraps it in `handlers/mcp.go` (clients, library, tool groups), `handlers/mcpserver.go`, `handlers/mcpsessions.go` and the OAuth2 set (`mcpoauth2.go`, `mcpoauth2discovery.go`, `mcpoauth2issuance.go`, `mcpoauth2consent.go`, `mcpoauth2sessions.go`, `mcpoauth2jwt.go`).

4.4 Key selection and networking

- `core/keyselectors` - weighted/random selection across multiple keys per provider.
- `core/network` - the outbound HTTP client for provider calls (proxy, TLS, timeouts, retries), driven by each provider's `NetworkConfig`.

5. Framework module (`framework/`)

5.1 ConfigStore

`framework/configstore` is the configuration database layer.

- Drivers: `postgres.go`, `mysql.go`, `sqlite.go`, `rdb.go`; interface in `store.go`; errors in `errors.go`.
- `migrations.go` plus `framework/migrator` (`addcolumn.go`, `dropcolumn.go`) handle schema evolution.
- `tables/` holds one file per domain: `user.go`, `workspace.go`, `enterprise.go`, `team.go`,

`customer.go`, `virtualkey.go`, `key.go`, `provider.go`, `model.go`, `modelconfig.go`,
`modelpricing.go`, `modelparameters.go`, `pricingoverride.go`, `budget.go`, `ratelimit.go`,
`routingrules.go`, `mcp*.go`, `mcpoauth2*.go`, `prompts.go`, `promptVersions.go`, `promptSessions.go`,
`skills.go`, `folders.go`, `featureflag.go`, `plugin.go`, `config.go`, `clientconfig.go`, `env.go`,
`sessions.go`, `temptokens.go`, `dlock.go` (distributed locks), `logstore.go`, `vectorstore.go`,
`smtplib.go`, `confighash.go`.

- `workspace.go` seeds the **RBAC catalog**, the three **system roles**, and exposes the resource/operation/name catalogs used by backend and UI.
- Secrets in the DB are encrypted with `framework/encrypt` when `UNIFAI_ENCRYPTION_KEY` is set (`tables/encryption.go`, `configstore/encryption.go`).

5.2 LogsStore

`framework/logstore` stores and aggregates request logs.

- Files: `logger.go`, `store.go`, `rdb.go`, `postgres.go`, `mysql.go`, `sqlite.go`, `hybrid.go`.
- `browser_ai.go` is the **entire Browser AI persistence layer** (see §9).
- Aggregations for dashboard/log pages: `matviews.go` (materialized views), `tables.go`, `cleaner.go` (retention), `asyncjob.go` (async inference jobs), `payload.go`.

5.3 RBAC (two enforcement layers, both applied to `/api/*`)

1. **Resource × Operation** - `framework/rbac/enforce.go`, `PathRequirementFor(method, path)`

returns the requirement for a route (or `nil` = not RBAC-gated). `admin` resolves to *allow all*;

an unknown/unseeded role yields an **empty set = deny**. **View** and **Read** are interchangeable.

2. **Sidebar section grants** - `framework/rbac/sections.go`, `SectionRequirementFor(method, path)`,

used for section-scoped sessions. Keys are `parent` or `parent/child`; `parent/*` accepts the parent or any child. Legacy `observability/browser-ai` is normalised to `browser-ai`.

Resource catalog (33): `GuardrailsConfig, GuardrailsProviders, GuardrailRules, UserProvisioning, Cluster, Settings, Users, Logs, Observability, Dashboard, VirtualKeys, ModelProvider, Plugins, MCPGateway, MCPToolGroups, MCPLogs, AdaptiveRouter, AuditLogs, Customers, Teams, RBAC, Governance, RoutingRules, PromptRepository, PromptDeploymentStrategy, SkillsRepository, AccessProfiles, APIKeys, Inference, Metrics, FeatureFlags, CircuitBreaker`

Operation catalog (6): **Read, View, Create, Update, Delete, Download**

System roles (auto-seeded in `configstore/workspace.go`, then merged on every boot so new permissions are picked up - except for roles a customer has customised):

Role	DAC	Grants
<code>admin</code>	<code>all-data</code>	every permission

Role	DAC	Grants
sub_admin	all-data	Read/View on Dashboard, Inference, Observability, MCPLogs, RoutingRules, Metrics ; full (Read/View/Create/Update/Delete/Download) on Logs - this is what gates all Browser AI write APIs; full on VirtualKeys, PromptRepository ; View/Read/Update on Governance ; View/Read/Create/Update on MCPGateway, MCPToolGroups ; View/Read on ModelProvider . No master provider keys, no system settings
user	own-data	Read/View only , on Dashboard, Logs, Inference, PromptRepository, Observability, MCPGateway (the readIDs set). The UI labels this "Prompt Repository only"; the backend also allows read-only views of logs/observability/MCP. Use allowed_sections to narrow what a user actually sees

Custom roles are created in the UI. **admin** is always re-synced to the full catalog.

*Browser AI note: every **/api/browser-ai/*** route maps to the **Logs** RBAC resource*

*(**PathRequirementFor**), so a role needs **Logs: View/Update** (or the **browser-ai** section grant)*

to see the Browser AI pages.

5.4 Other framework packages

Package	Purpose	API surface
cluster	multi-node config propagation	/api/cluster, /internal/cluster/kv
loadbalancer	adaptive routing across providers/keys	/api/load-balancer, /api/load-balancer/routes
routing	routing-rule engine	/api/governance/routing-rules, /api/routing-rules
circuitbreaker	failure isolation policies	/api/circuit-breaker/policies, /state
featureflags	server-side flags	/api/feature-flags
rbac, oauth2, temptoken	permissions, OAuth2 provider, short-lived tokens	/api/rbac/*, /api/scim/oauth/*
mailer, alerts	SMTP and alert delivery (webhooks)	/api/smtp-config, /api/alert-channels
connectors	external sinks (BigQuery, Kafka, Datadog, New Relic, PubSub)	/api/connectors
objectstore, vectorstore, kvstore	blob / vector / KV backends (Qdrant, Weaviate, Redis)	/api/vector-store-config, /api/cache

Package	Purpose	API surface
tracing	trace store + spans	/api/logs trace views
querscope, mcpcatalog, mcptoolgroups, mcp_headers, modelcatalog, encrypt, envutils	supporting services	-
streaming	SSE/stream utilities shared by transports	-

6. Plugins (**plugins/**)

A plugin implements one or more core interfaces and is loaded either as a **built-in** (by name) or

as a **dynamic .so** via **SharedObjectPluginLoader** (**server/plugins.go**).

Built-in load order (fixed placement indexes):

#	Plugin	Role
1	telemetry	request metrics/counters
2	prompts	prompt repository integration
3	logging	LLM/MCP log persistence
4	governance	virtual keys, budgets, rate limits, teams, customers, business units
5	otel	OpenTelemetry export
6	semanticcache	embedding-based response caching
7	compat	parameter compatibility/conversion between client and provider shapes
8	maxim	Maxim observability integration
-	guardrails	rule-based input/output validation + provider config
-	connectors	external log/analytics sinks
-	modelcatalogresolver	resolves model metadata/pricing
-	mock	mock provider for tests/E2E
-	jsonparser	tolerant JSON parsing helpers

Plugin config lives in **config.json** under **plugins**; plugins can be enabled/disabled and reordered at runtime (**GET/POST /api/plugins**). A plugin that fails to load is marked disabled and reported in the startup plugin-status table (**plugin status: <name> - <status>**).

7. HTTP surface (**transports/unifai-http**)

7.1 Entry point and flags

main.go:

Flag	Env fallback	Default
-port	APP_PORT	8001
-host	APP_HOST / UNIFAI_HOST	localhost
-app-dir	APP_DIR	./data
-log-level	LOG_LEVEL	info
-log-style	LOG_STYLE	json

Startup prints an ASCII banner (**-ldflags -X main.Version=...**), starts the pprof server, configures the logger, then **Bootstrap()** -> **Start()**.

7.2 Middleware assembly

Server handler wrapper (**server.go**):

```
SecurityHeadersMiddleware → CORS → RequestDecompression → router
```

API middleware chain (**apiMiddlewares**, applied to every **/api/*** route):

```
AuthMiddleware.APIMiddleware() session cookie / Bearer session token WorkspaceAuditMiddleware audits mutating /api/* into Audit Logs RBACMiddleware resource × operation + section grants
```

Inference middleware chain (**/v1/*** and pass-through):

```
TracingMiddleware → TransportInterceptor → AuthMiddleware.InferenceMiddleware → handler
```

TransportInterceptor runs **HTTPTransportPostHook** plugins and captures their log entries; the tracing defer runs afterwards so plugin logs land inside the trace.

7.3 Request identification on inference

- **Authorization: Bearer <virtual-key|provider-key>** or **x-api-key: <key>**.
- Dashboard-originated inference is enriched from the session cookie

(**enrichInferenceFromDashboardSession**); provider headers are normalised to a single

Authorization: Bearer ... (**x-api-key** is deleted).

- Password-as-Bearer (base64 **user:pass**) is **removed** - only session tokens are accepted as Bearer.
- Virtual Key self-service quota: **GET /api/governance/virtual-keys/quota** (the VK header is the credential).

7.4 Route groups

Group	Representative routes	Handler file(s)
Health / UI	GET /health, GET /robots.txt, / (SPA)	health.go, ui.go
Session / auth	/api/session/login, /logout, /register, /register/verify, /register/resend, /forgot-password, /verify-otp, /reset-password, /forgot-username, /is-auth-enabled, /ws-ticket, /api/session/users*	session.go, session_register_verify.go, auth_security.go

Group	Representative routes	Handler file(s)
Workspace / RBAC	<code>/api/users/{id}/role, /api/roles*, /api/permissions, /api/resources, /api/operations, /api/rbac/*, /api/access-profiles*</code>	<code>workspace.go, rbac.go, rbac_enforce.go, accessprofiles.go</code>
Governance	<code>/api/governance/virtual-keys*, /customers*, /teams*, /business-units*, /routing-rules*, /complexity-analyzer-config, /api/guardrails/*</code>	<code>governance.go, businessunits.go, team_members.go, guardrails.go</code>
Inference	<code>/v1/chat/completions, /v1/completions, /v1/responses*, /v1/embeddings, /v1/rerank, /v1/ocr, /v1/audio/*, /v1/images/*, /v1/videos*, /v1/batches*, /v1/files*, /v1/containers*, /v1/realtime, /v1/async/*</code>	<code>inference.go, asyncinference.go, realtime_*.go, webrtc_realtime.go</code>
Pass-through	<code>/openai/*, /anthropic/*, /genai/*, /bedrock/*, /cohere/*, /cursor/*, /litellm/*, /langchain/*, /pydanticai/*</code>	<code>integrations/*.go</code>
Observability	<code>/api/logs*, /api/logs/dashboard, /api/logs/filterdata, /api/logs/rankings*, /api/logs/histogram/*, /api/mcp-logs*, /api/audit-logs, /api/connectors, /api/alert-channels</code>	<code>logging.go, auditlogs.go, alertchannels.go, integrations.go</code>
Providers / models	<code>/api/providers*, /api/keys*, /api/models*, /api/model-configs*</code>	<code>providers.go, provider_keys.go, list_models_vk.go</code>
MCP	<code>/api/mcp/clients, /api/mcp/library*, /api/mcp/client*, /api/mcp/tool-groups*, /api/mcp-logs, /api/mcp/inference</code>	<code>mcp.go, mcpserver.go, mcpsessions.go, mcptoolgroups.go, mcphheaders.go</code>
Prompts / skills	<code>/api/prompt-repo*, /api/prompt-deployments, /api/skills*, /api/skills/serve/*</code>	<code>prompts.go, prompts_lifecycle.go, skills.go, skills_serving.go</code>
Config / platform	<code>/api/config, /api/proxy-config, /api/vector-store-config, /api/smtp-config, /api/cache, /api/plugins, /api/feature-flags, /api/cluster, /api/circuit-breaker/*, /api/load-balancer*</code>	<code>config.go, plugins.go, featureflags.go, cluster.go, circuitbreaker.go, cache.go</code>
Enterprise / SCIM	<code>/api/scim/*, /api/scim/oauth/*</code>	<code>scim.go, scim_handlers.go, scim_oauth.go, scim_auth.go</code>
Browser AI	see §9	<code>browser_ai*.go</code>
WebSocket	<code>/ws, WS tickets, responses/realtime WS</code>	<code>websocket.go, wsrealtime.go, wsresponses.go, ws_ticket.go</code>
Dev	<code>/api/dev/* + pprof (dev mode only)</code>	<code>devpprof*.go</code>

7.5 Health and observability endpoints

- **GET /health** - pings ConfigStore, LogsStore and VectorStore concurrently (10 s timeout) and returns `{"status": "ok", "components": {"db_pings": "ok"}}`; returns HTTP 503 with the first failing component otherwise. Can be short-circuited by `client.disable_db_pings_in_health`.
- **GET /robots.txt** - static **Allow: /**.
- **GET /api/version** - build version.
- Logging skip list (no request logging): `/health`, `/_next`, `/api/dev`.
- Metrics (Prometheus) and traces come from the telemetry/otel plugins and the Logs API.
- Dev builds additionally expose pprof handlers (gated by `devpprof_prod.go` in release builds).

8. Data model (key tables)

8.1 ConfigStore tables

- **governance_users** - `id, username(uniq), email, password, role(admin|sub\_admin|user), status (pending|approved|rejected|email\_unverified|disabled), budget, rate\_limit, budget\_id, rate\_limit\_id, allowed\_prompt\_repos, allowed\_sections, reviewed\_at, external\_id, created\_at, updated\_at`. **IsApproved()** treats an empty status as approved (legacy rows).
- **rbac_roles** - `id, name(uniq), description, is\_system\_role, dac(all-data|own-data), permission\_ids(JSON), created\_at, updated\_at`.
- **access_profiles** - `id, name(uniq), description, is\_active, version, calendar\_aligned, tags(JSON), spec(JSON), created\_at, updated\_at` (reusable governance policy templates with activate / deactivate / clone operations).
- **governance_business_units** - `id, name(uniq), team_ids(JSON), budget(JSON), rate_limit(JSON)`.
- Grouped in `tables/workspace.go`: `alert_channels`, `circuit_breaker_policies`, `mcp_tool_groups`, `prompt_deployments`, `workspace_settings`, `virtual_key_users`, `virtual_key_teams`, `virtual_key_customers`, `audit_logs`.
- Plus the core governance tables: teams, customers, virtual keys, provider keys, models, model configs, pricing + pricing overrides, budgets, rate limits, routing rules, prompts/versions/sessions, skills, folders, feature flags, plugins, sessions, temp tokens, distributed locks, SMTP, vector-store config, MCP clients/library/OAuth2 issuance.

8.2 Browser AI tables (`framework/logstore/browser_ai.go`)

Table	Purpose	Notable columns
browser_ai_logs	one row per intercepted prompt / upload	id, timestamp, platform, domain(host), user_prompt_full, action, status, rule_triggered, risk_score, predictive_risk, agent_id, agent_hostname, client_ip, reply_bot_text, attachment_name/type/size/path/expiry, metadata(JSON)
browser_ai_search_logs	search-engine queries + clicked results	id, timestamp, engine, browser, is_incognito, query, clicked_url, clicked_title, url, host, agent_id, predictive_risk
browser_guard_rules	guard rules	id, name, pattern, rule_type(regex or ai_bot), severity, action, warning_message, active, description, bot_provider, bot_model, bot_prompt, bot_reference_image, bot_reference_image_type
browser_ai_target_websites	monitored AI sites	id, domain(uniq), platform_name, monitored, block_entire_website, host_role(ui/chat/file), status(MONITORED/PAUSED/BLOCKED), parent_id
browser_ai_control_settings	single-row interaction policy (browser-controls-default)	upload/download warning text and policy toggles
browser_ai_agents	Guard fleet inventory	id, hostname, agent_type(endpoint/network), status(active/uninstall_pending/uninstalled), last_seen_at, health_detail, contact_email(+pinned), uninstall_key_hash, uninstall_key_enc, has_uninstall_key, uninstall_key_rotated_at, proxy_bundle_sha
browser_ai_agent_settings	company-level uninstall key	id, uninstall_key_enc
browser_guard_fleet_config	company Guard defaults that agents pull on heartbeat	id, ... notes
browser_guard_rebuild_logs	Guard package rebuild history	id, status(success/failed), bundle_sha, message, log, created_at
browser_ai_warning_emails	employee warning e-mail log	id, sent_at, to_email, message, agent_id, error_detail

- Attachment bytes stay on disk (**APP_DIR/attachments**); the DB stores size, path and expiry metadata.

- **BrowserAIDManager** (same file) implements every CRUD/aggregation API used by the handlers,

including **GuardSeverityScore** (severity -> predictive risk score + label) and

NormalizeBrowserAIAgentType.

9. Browser AI (UnifAI Guard) - technical detail

9.1 Backend routes (**handlers/browser_ai.go** registers all of these)

Logs / insights

```
GET /api/browser-ai/logs list + filter prompt logs GET /api/browser-ai/logs/stats aggregated stats DELETE
/api/browser-ai/logs delete all POST /api/browser-ai/logs/bulk-delete DELETE /api/browser-ai/logs/{id} GET
/api/browser-ai/search-logs POST /api/browser-ai/search-logs agent ingest DELETE /api/browser-ai/search-logs
POST /api/browser-ai/search-logs/bulk-delete DELETE /api/browser-ai/search-logs/{id} GET
/api/browser-ai/insights/stats Guard Insights
```

Rules / controls / targets

```
GET|POST /api/browser-ai/rules + POST /rules/import, PUT|DELETE /rules/{id} POST
/api/browser-ai/rules/test-bot test an AI Guard Bot rule POST /api/browser-ai/rules/generate-regex generate
regex from a policy description GET|PUT /api/browser-ai/controls interaction controls (single row) GET
/api/browser-ai/targets + POST /targets, POST /targets/import, PUT|DELETE /targets/{id}
```

Fleet / agents / setup

```
GET /api/browser-ai/agents list agents (filters: status, source, search) POST
/api/browser-ai/agents/heartbeat agent registration + health (Guard secret) POST
/api/browser-ai/agents/bulk-delete DELETE /api/browser-ai/agents/{id} PUT
/api/browser-ai/agents/{id}/contact-email GET /api/browser-ai/agents/settings GET|PUT
/api/browser-ai/agents/uninstall-key company uninstall key GET /api/browser-ai/agents/{id}/uninstall-key
POST /api/browser-ai/agents/{id}/uninstall-key/rotate POST /api/browser-ai/agents/uninstall-verify employee
key check POST /api/browser-ai/agents/uninstall agent requests removal POST
/api/browser-ai/agents/uninstall-ack POST /api/browser-ai/agents/uninstall-status POST
/api/browser-ai/agents/{id}/remote-uninstall admin-initiated removal GET|PUT /api/browser-ai/fleet-config
company-wide Guard defaults GET /api/browser-ai/setup/info GET /api/browser-ai/setup/download.zip aliases:
-windows.zip, -mac.zip POST /api/browser-ai/setup/rebuild rebuild + republish Guard packages GET
/api/browser-ai/setup/rebuild-history GET /api/browser-ai/setup/proxy-bundle.json GET
/api/browser-ai/setup/proxy-bundle.zip hot-update bundle for installed Guards POST
/api/browser-ai/send-warning-email
```

Ingest / PAC / helpers

```
POST /api/browser-ai/intercept agent prompt/upload event POST /api/browser-ai/intercept-file agent file
event GET /api/browser-ai/attachments/{id} attachment download (auth-gated) GET /api/browser-ai/proxy.pac
PAC script GET /api/browser-ai/pac PAC script (alias used by GPO / network proxy) GET
/api/browser-ai/ollama-models models available for AI Guard Bot rules
```

9.2 Guard-fleet authentication (no user session)

isPublicBrowserAIRoute and **isGuardKeyBrowserAIRoute** (**handlers/middlewares.go**) let the Guard agent call the backend **without a dashboard session cookie**:

- **Public (secret-validated):** **intercept**, **intercept-file**, **search-logs** (POST), **proxy.pac**, **pac**, **agents/heartbeat**, **agents/uninstall-verify**, **agents/uninstall**, **agents/uninstall-ack**, **agents/uninstall-status**.
- **Guard-key:** **targets**, **rules**, **controls**, **fleet-config** (GET), **setup/download*.zip**, **setup/proxy-bundle.json**, **setup/proxy-bundle.zip**.
- The fleet secret is **UNIFAI_GUARD_SECRET**. With **UNIFAI_GUARD_REQUIRE_SECRET=1** (production default

in `.env.example` and compose) agent APIs **fail closed** when the secret is missing or wrong (`browser_ai_guard_security.go`).

- `UNIFAI_PAC_ALLOW_QUERY_PROXY=0` (default) makes the PAC ignore a client-supplied `?proxy=` and use only fleet/env values - anti PAC-hijack.

9.3 Guard desktop agent (`apps/browser-guard/agent`)

`unifai_agent.py` is the entry point. Supporting modules:

Module	Responsibility
<code>guard_bootstrap.py</code>	single-instance lock, logging bootstrap, runtime version resolution
<code>agent_config.py</code>	env -> JSON config precedence (<code>UNIFAI_BACKEND_URL</code> , <code>UNIFAI_PROXY_ADDR</code> , <code>UNIFAI_PAC_URL</code> , <code>UNIFAI_PAC_SYNC_SECONDS</code> , <code>UNIFAI_SERVER_MODE</code> , <code>UNIFAI_AGENT_TYPE</code> , <code>UNIFAI_LISTEN_HOST</code> , <code>UNIFAI_PAC_ADVERTISE_ADDR</code> , <code>UNIFAI_AGENT_ID</code> , <code>UNIFAI_AGENT_HOSTNAME</code>)
<code>agent_identity.py</code>	persistent agent ID + host/OS metadata for the fleet inventory
<code>agent_heartbeat.py</code>	periodic heartbeat, fleet config pull, remote-uninstall intent, health detail
<code>agent_health.py</code>	port checks (<code>port_open</code> , <code>free_proxy_port</code>), health loop
<code>agent_pac_content.py</code>	backend reachability check, PAC fetch, local PAC write
<code>agent_pac_server.py</code>	local PAC/health HTTP server (<code>PAC_HTTP_PORT</code> , default 18195, falls back to +1/+2 if busy)
<code>agent_pac_orchestration.py</code>	apply PAC with cache-bust, clear/restore runtime, fail-open direct mode
<code>agent_proxy_engine.py</code>	runs the mitmproxy worker child process
<code>agent_proxy_bundle.py</code>	SHA-tracked hot-update bundle, mark-bad, relaunch, running SHA
<code>agent_certs.py</code>	installs the mitmproxy CA into the OS trust store, writes <code>ca_install_status.txt</code>
<code>agent_browser_policy.py</code>	browser QUIC / DoH policy adjustments
<code>agent_autoupdate.py</code>	compares <code>UNIFAI_GUARD_RUNTIME_VERSION</code> with the installer version and updates
<code>agent_lifecycle.py</code>	first-run prompt, uninstall flow (key-gated), user messages
<code>agent_logging.py</code>	file logging, <code>get_resource_path</code>
<code>agent_http.py</code> , <code>agent_state.py</code> , <code>guard_platform.py</code>	HTTP client, local state, OS paths (<code>data_dir</code> , <code>log_hint_path</code> , <code>register_autostart</code>)

Default paths: data `%LOCALAPPDATA%\UnifAI\Guard` (Windows) /

`~/Library/Application Support/UnifAI/Guard` (macOS); log `unifai_guard.log`;

CA status `ca_install_status.txt`.

9.4 mitmproxy addon (`apps/browser-guard/proxy/browser_ai_proxy.py`)

The addon is **not** a normal Python package. `unifai_proxy_parts/MANIFEST.txt` lists the parts that are loaded **in order into one shared namespace** (preserving the original single-file monolith behaviour):

```
config_caches_rules.py config, caches, target/rule fetch + matching helpers_prompts.py prompt extraction +
helper utilities uploads_detect.py upload / file detection file_policy.py per-file-type policy decisions
extract_office_backend.py Office/PDF/image text extraction + backend calls responses_inject.py response
rewriting / block injection (SSE aware) responses_addon.py response learning (e.g. filename discovery)
```

Flow: target-domain match -> extract prompt/attachment -> evaluate rules (regex locally; **ai_bot** via the configured LLM/Ollama) -> verdict **Allowed | Redacted | Warned | Blocked | SiteBlocked | Bot Answered** -> inject the browser-visible response (JSON error, Anthropic SSE **unifai_guard_blocked**, etc.) -> POST the event to **/api/browser-ai/intercept**.

Editing any part requires rebuilding the Guard (**installer/build_installer.bat** on Windows, **installer/build_macos.sh** on macOS) and keeping cross-part names intact.

9.5 Guard distribution and hot update

- Build: **make build-guard-windows** (Inno Setup) / **make build-guard-mac** (+ **-pkg**) - see **apps/browser-guard/README.md**. Without **make**, run **installer\build_installer.bat**.
- Config sync: **python apps/browser-guard/scripts/sync_config_from_env.py** regenerates **config/unifai_guard_config.json** and the docs from **.env** (**SERVER_DOMAIN**, **UNIFAI_PROXY_ADDR**, **PAC_HTTP_PORT**, **UNIFAI_GUARD_SECRET**).
- Server-side artifacts live in **apps/browser-guard/release/** and are copied into the image (**/app/release**), so **Browser AI -> Setup -> Download Setup ZIP** serves them.
- **Browser AI -> Setup -> Rebuild & Publish** re-publishes the proxy/agent Python sources as a bundle (**/api/browser-ai/setup/proxy-bundle.zip**); installed agents pick it up via **agent_proxy_bundle.py**.
- Version file: **apps/browser-guard/release/VERSION.txt** -> **1.1.15**.

10. Frontend (ui/)

10.1 Build and integration

- Vite + React 19 + TanStack Router (file-based routes under **ui/app/****, generated into **ui/app/routeTree.gen.ts**).
- State: Redux Toolkit with RTK Query APIs under **ui/app/_fallbacks/enterprise/lib/store/apis/*** (**accessProfileApi**, **alertChannelsApi**, **auditLogsApi**, **businessUnitsApi**, **circuitBreakerApi**, **clusterApi**, **connectorsApi**, **loadBalancerApi**, **mcpToolGroupsApi**, **promptDeploymentsApi**, **rbacApi**, **scimApi**, **virtualKeyUsersApi**, ...).
- **npm run build:unifai** = rebrand + **vite build** + typecheck; **npm run copy-build** copies **out/** into **transports/unifai-http/ui**, which the Go binary embeds (**//go:embed**). The UI therefore ships **inside** the server binary - no separate web server is required.
- Dev server: **npm run dev**; **make dev** runs UI + API together.
- Quality gates: **npm run typecheck** (tsc --noEmit), **npm run lint** (oxlint), **npm run format** (oxfmt).

10.2 Pages

Top level: `/` (landing), `/login`, `/signup`, `/oauth/consent`, `/pprof`, `/workspace/**`, plus `__root`, `__notFound`, `__error`.

Workspace pages (one `page.tsx` per route, each with a `layout.tsx`):

Section	Routes
Observability	<code>dashboard</code> , <code>logs</code> , <code>logs/connectors</code> , <code>mcp-logs</code> , <code>observability</code> , <code>audit-logs</code> , <code>alert-channels</code>
Browser AI	<code>browser-ai</code> (tabs: <code>overview</code> , <code>logs</code> , <code>search-logs</code> , <code>rules</code> , <code>targets</code> , <code>agents</code> , <code>telemetry</code> , <code>setup</code>)
Models	<code>providers</code> , <code>providers/model-limits</code> , <code>providers/routing-rules</code> , <code>model-catalog</code> , <code>model-limits</code> , <code>custom-pricing</code> , <code>custom-pricing/overrides</code> , <code>complexity-router</code> , <code>routing-rules</code> , <code>routing-rules/tree</code> , <code>circuit-breaker</code>
MCP Gateway	<code>mcp-gateway</code> , <code>mcp-registry</code> , <code>mcp-registry/library</code> , <code>mcp-registry/oauth-callback</code> , <code>mcp-tool-groups</code> , <code>mcp-sessions</code> (+ <code>auth</code> , <code>auth-failed</code> , <code>auth-success</code>), <code>oauth-grants</code> , <code>mcp-auth-config</code> , <code>mcp-settings</code>
Governance	<code>governance</code> , <code>governance/virtual-keys</code> , <code>governance/users</code> , <code>governance/teams</code> , <code>governance/business-units</code> , <code>governance/customers</code> , <code>governance/rbac</code> , <code>governance/access-profiles</code> , <code>rbac</code> , <code>virtual-keys</code> , <code>scim</code> , <code>scim/oauth-discover-callback</code>
Guardrails	<code>guardrails</code> , <code>guardrails/configuration</code> , <code>guardrails/providers</code>
Platform	<code>plugins</code> , <code>cluster</code> , <code>adaptive-routing</code> (+ <code>settings</code>), <code>prompt-repo</code> (+ <code>prompts</code>), <code>skills-repo</code> , <code>docs</code> , <code>config</code> (+ <code>api-keys</code> , <code>caching</code> , <code>client-settings</code> , <code>compatibility</code> , <code>feature-flags</code> , <code>logging</code> , <code>mcp-gateway</code> , <code>observability</code> , <code>performance-tuning</code> , <code>pricing-config</code> , <code>proxy</code> , <code>security</code>)

10.3 Sidebar and section gating

`ui/components/sidebar.tsx` builds the menu from `ui/lib/constants/workspaceSections.ts`:

```
observability | browser-ai | models | mcp-gateway | plugins | governance | guardrails | cluster-config |
adaptive-routing | prompt-repository | skills-repository | settings
```

Each section has a `defaultPath` and optional child `items`. Visibility is decided by:

1. RBAC hooks - `useRbac(RbacResource.X, RbacOperation.View)`.
2. Role checks for admin-only pages (Users requires `admin` or `sub_admin`).
3. Section grants for section-scoped sessions (`scopedSidebarSections` from `authStatus.allowed_sections`), with automatic redirect to `getDefaultPathForSections(...)`

when the user lands on a route they are not granted.

Browser AI is reachable with `Logs` view RBAC or the `admin/sub_admin` role

(`hasBrowserAiAccess`), matching the backend mapping of `/api/browser-ai/*` -> `Logs`.

10.4 Browser AI UI implementation files

ui/app/workspace/browser-ai/: **page.tsx** (tabs, tables, dialogs), **browserAiTypes.ts**, **browserAiConstants.ts** (**GUARD_BOT_OLLAMA_PROVIDER** = "ollama", **GUARD_BOT_OLLAMA_MODEL** = "llama3.2", reference image max 512 KB), **browserAiFormat.ts**, **browserAiLogHelpers.ts**, **browserAiGuardHelpers.ts**, **guardBotModelPickers.tsx**, **guardRuleAIEvaluatorFields.tsx**, **guardRuleImportDialog.tsx**, **targetImportDialog.tsx**, **regexLiveTestPanel.tsx**, **logPromptPreviewCell.tsx**, **logTimestampCell.tsx**, **logBadges.tsx**, **attachmentPreview.ts**, **relatedHosts.ts**, plus import templates under **examples/**.

Types as implemented: **GuardRuleAction** = "BLOCK" | "REDACT" | "WARN", **GuardRuleSeverity** = "CRITICAL" | "HIGH" | "MEDIUM", **GuardRuleEvalMode** = "ai" | "regex".

11. Security model

Area	Implementation
Dashboard sessions	Session rows stored in DB; token cookie with HttpOnly + SameSite=Lax (+ Secure only for real TLS or a trusted proxy sending X-Forwarded-Proto: https); default 24 h TTL
WebSocket auth	Cookie-only (/api/session/ws-ticket issues short-lived tickets); legacy ?token= query auth removed (token-leak risk)
Inference auth	Virtual Key or provider key via Authorization: Bearer / x-api-key ; enforce_auth_on_inference flag
Brute force	Per-IP: 30 failed logins / 15 min window; per-username lockout rows in the DB; remaining minutes returned to the client
Password reset	PASSWORD_RESET_SECRET (min 32 chars) HMAC-signed token, 10 min expiry; 5 requests/hour/target + 10/hour/IP; OTP max 5 attempts; 2 min cooldown
Registration	Self-signup with e-mailed code (email_unverified) -> admin approval (pending -> approved) or rejected ; SCIM can set disabled
Authorization	Two RBAC layers (resource×operation, section grants) + audit middleware recording mutating /api/* actions
Secrets at rest	framework/encrypt encryption for DB-stored secrets when UNIFAI_ENCRYPTION_KEY is set (provider keys, alert channels, uninstall keys)
Guard fleet	Shared UNIFAI_GUARD_SECRET , fail-closed with UNIFAI_GUARD_REQUIRE_SECRET=1 ; PAC hijack protection via UNIFAI_PAC_ALLOW_QUERY_PROXY=0
Guard uninstall	Company and per-agent uninstall keys, stored hashed + encrypted, rotatable, with an audit trail and remote-uninstall flow
Transport	TLS terminated at the edge/reverse proxy; TRUST_PROXY_HEADERS=1 honours X-Forwarded-Proto/X-Forwarded-For only from trusted peers (TRUSTED_PROXIES required for public edges)

Area	Implementation
CORS	<code>client.allowed_origins</code> (from <code>SERVER_DOMAIN</code>); optional <code>ALLOW_LOCALHOST_CORS=1</code> for a local UI against a remote API
Headers	<code>SecurityHeadersMiddleware</code> applied to every response
Open bootstrap	When auth is disabled, unauthenticated admin is allowed only from loopback or with <code>ALLOW_OPEN_AUTH=1</code> (local bootstrap only)
Inference hardening	Provider headers normalised (<code>x-api-key</code> removed), password-as-Bearer removed, error sanitisation (<code>errorsanitizer.go</code>)
Destructive ops	Admin-only guards (<code>adminOnlyLogs</code> , <code>requireGuardAdmin</code>); attachment access privileged

12. Build, test and developer workflow

12.1 Gateway

```
make setup-workspace # go work use ./core ./framework ./plugins/... ./transports make dev # UI (vite) + API
with hot reload (air) make build # build UI then the Go binary make run # build + run make fmt ; make lint #
gofmt / linter make test # unifai-http transport tests (gotestsum) make test-core # core provider tests
(PROVIDER=openai TESTCASE=... PATTERN=...) make test-framework # framework tests make test-plugins # plugin tests
make test-http-transport # HTTP transport tests make test-governance # governance tests make test-mcp # MCP
tests (TYPE=... TESTCASE=...) make test-semantic-cache # semantic cache e2e make test-all # all of the above make
run-e2e ; make run-e2e-ui ; make run-e2e-headed ; make run-e2e-api make test-integrations-py ; make
test-integrations-ts make run-provider-harness-test ; make run-cli-harness-test
```

Reports: `make generate-html-reports`; cleanup with `make clean-test-reports`.

12.2 UI

```
cd ui npm ci npm run dev # Vite dev server npm run typecheck # tsc --noEmit npm run lint # oxlint npm run
build # vite build + typecheck + copy into transports/unifai-http/ui
```

12.3 Guard

```
make sync-guard-config # regenerate Guard config/docs from .env make build-guard-windows # Windows: Setup.exe
+ Windows ZIP (Inno Setup 6 required) make build-guard-mac # macOS: .app + macOS ZIP (must run on a Mac) make
build-guard-mac-pkg # macOS: .pkg from the existing .app make help-guard # list Guard targets
```

Guard unit tests live in `apps/browser-guard/agent/tests/` and `apps/browser-guard/proxy/tests/`
(PAC targets, proxy bundle, file-type extraction, latency, rule actions, target predict, upload names).

12.4 Docker / Helm

```
make docker-image # LOCAL=1 uses deploy/docker/Dockerfile.local (go workspace based) make docker-run # run
the container (CONFIG=... optional) make helm-index # repackage chart + regenerate deploy/helm/index.yaml
```

12.5 Conventions

- Go modules resolve through `go.work`; add new modules with `go work use`.
- Handlers keep the `RegisterRoutes(router, middlewares...)` pattern and are grouped by feature.

- UI files use camelCase names, Radix primitives from **components/ui**, and RTK Query for API calls.
- Guard proxy parts must stay in **MANIFEST.txt** order.

13. Extension points

Goal	Where to work
Add an LLM provider	new package under core/providers/<name> implementing the provider contract, then register it in the provider registry
Add a pass-through API surface	new file in transports/unifai-http/integrations/ + registration in integrations/router.go
Add a plugin	new module under plugins/<name> exposing PluginName ; add a case in loadBuiltinPlugin (server/plugins.go) and add the module to go.work
Add a dashboard API	new handlers/<feature>.go with RegisterRoutes , then call it from server.RegisterAPIRoutes
Gate a new route with RBAC	extend PathRequirementFor (framework/rbac/enforce.go) and, for section-scoped sessions, SectionRequirementFor (framework/rbac/sections.go)
Add an RBAC resource	add it to RBACResourceNames (framework/configstore/workspace.go); permissions are generated from the cartesian catalog
Add a UI page	create ui/app/workspace/<route>/page.tsx + layout.tsx , add the section item in ui/lib/constants/workspaceSections.ts , and the nav entry in ui/components/sidebar.tsx
Add a Guard detection	edit the relevant part in apps/browser-guard/proxy/unifai_proxy_parts/ (keep MANIFEST order) and rebuild the Guard
Add a Guard rule type	BrowserGuardRule (framework/logstore/browser_ai.go) + the UI rule editor + the proxy evaluator
Change gateway guardrail behaviour	plugins/guardrails + /api/guardrails/*

14. Glossary

Term	Meaning in this codebase
Virtual Key (VK)	A gateway-issued credential that maps to budgets/limits/teams/customers and is used by inference clients
ConfigStore	Config/workspace database (users, roles, keys, providers, settings)
LogsStore	Log database (LLM logs, MCP logs, Browser AI logs, aggregates)
Section grant	Sidebar-section permission (governance/users) used for section-scoped sessions
DAC	Data access scope on a role: all-data or own-data
Target website	A monitored AI domain (with platform name, host role and monitored/blocked flags)

Term	Meaning in this codebase
Guard rule	A detection rule: regex (pattern) or ai_bot (LLM-judged, optional reference image)
Guard agent	A Guard installation reporting to the fleet (endpoint or network)
Fleet config	Company-wide Guard defaults that agents pull on heartbeat
PAC	Proxy auto-config script served by the backend, pointing browsers at the Guard proxy
Proxy bundle	Server-published hot-update bundle of the Guard proxy/agent Python code
Uninstall key	Company/per-agent secret required to turn the Guard off
Intercept	The Guard -> backend event that records a prompt, upload or verdict