

UnifAI / Raksha Enterprise Administrator Guide

Complete Admin Console Manual: Configuration, Governance & Fleet Management

Document Version: 2.4.0

Classification: Enterprise Administration & Security Governance

Target Audience: Platform Administrators, InfoSec Teams, Compliance Officers, Workspace Managers

1. Initial Access & First-Run Security Checklist

Upon launching the UnifAI platform, follow this mandatory security onboarding sequence before inviting users or onboarding endpoints:

1. Initial Root Sign-In:

- Navigate to your instance: <https://unifai.yourcompany.com/auth/login>.
- Authenticate using the bootstrap credentials defined in your server environment (**ADMIN_EMAIL** and **ADMIN_PASSWORD**).

2. Immediate Password Rotation:

- Click your profile avatar in the top-right corner \rightarrow **Account Settings** \rightarrow **Security**.
- Change the password immediately to a strong, high-entropy passphrase (minimum 16 characters).

3. Dedicated Admin Account Provisioning:

- Navigate to **Governance** \rightarrow **Users**.
- Create a named personal administrator account with the **admin** role.
- Demote or securely store the bootstrap account for break-glass emergency use only.

4. SMTP / Email Notification Setup:

- Navigate to **Settings** \rightarrow **SMTP Configuration**.
- Provide your corporate SMTP gateway credentials (host, port 587/465, TLS, user, password, from-address).
- Send a test email to verify that registration verification codes, password resets, and high-severity compliance alerts function correctly.

5. Branding & Legal Notice Customization:

- Under **Settings** \rightarrow **General Settings**, update the Company Name, Portal Logo, and End-User Terms of Use notice shown upon login.
-

2. Governance, User Lifecycle & RBAC Administration

2.1 User Lifecycle Management

Located at **Governance** \rightarrow **Users** (</workspace/governance/users>):

Operation	Steps in Admin UI	Underlying API Route
Approve Pending User	Select user row in Pending tab \rightrightarrow Click Approve . Assign Role and initial Allowed Sections.	POST /api/session/users/{id}/approve
Reject Sign-Up	Select user row \rightrightarrow Click Reject . Enter optional explanation reason.	POST /api/session/users/{id}/reject
Direct User Creation	Click + Create User . Enter Email, Full Name, Initial Password, Role, Budget, and Rate Limit.	POST /api/session/users
Role Assignment	Open user profile \rightrightarrow Select Role dropdown (admin , sub_admin , user , or Custom Role).	PUT /api/users/{id}/role
Section-Scoped Access	Under Allowed Sections , toggle specific sidebar items (e.g. Prompt Repository, Logs, Browser AI).	PUT /api/session/users/{id}
Account Deactivation	Click user actions menu \rightrightarrow Disable Account . Immediately revokes all active sessions.	DELETE /api/session/users/{id}

2.2 Roles & Permissions Engine (RBAC)

Located at **Governance** \rightrightarrow **Roles & Permissions** (**/workspace/governance/roles**):

UnifAI features a granular enterprise authorization matrix based on:

- **33 System Resources:** **GuardrailsConfig**, **GuardrailsProviders**, **GuardrailRules**, **UserProvisioning**, **Cluster**, **Settings**, **Users**, **Logs**, **Observability**, **Dashboard**, **VirtualKeys**, **ModelProvider**, **Plugins**, **MCPGateway**, **MCPToolGroups**, **MCPLogs**, **AdaptiveRouter**, **AuditLogs**, **Customers**, **Teams**, **RBAC**, **Governance**, **RoutingRules**, **PromptRepository**, **PromptDeploymentStrategy**, **SkillsRepository**, **AccessProfiles**, **APIKeys**, **Inference**, **Metrics**, **FeatureFlags**, **CircuitBreaker**.
- **6 Discrete Operations:** **Create**, **Read**, **Update**, **Delete**, **View**, **Download**.
- **Data Access Control (DAC):**
 - **all-data:** Role grants visibility across the entire enterprise tenant.
 - **own-data:** Restricts user visibility strictly to records, keys, or logs authored by the authenticated user.

Creating a Custom Role:

1. Click + **Create Custom Role**.
2. Name the role (e.g., **ComplianceAuditor**, **PromptEngineerLead**).
3. Set the Data Access Control scope (**all-data** or **own-data**).
4. Select the specific resource checkboxes and assign allowable operations.
5. Click **Save Role**. Custom roles immediately appear in user assignment menus.

2.3 Access Profiles (Policy Templates)

Located at **Governance** \rightrightarrow **Access Profiles** (**/workspace/governance/access-profiles**):

Access Profiles allow you to bundle permissions, model allowances, rate limits, and allowed sections into reusable, versioned policy packages.

- **Creating a Profile:** Click + **New Access Profile**, enter Name (e.g., **Contractor-Limited-Profile**), Description, and assign specific AI models allowed for inference.

- **Binding to Users:** When onboarding dozens of employees, simply select the Access Profile to automatically synchronize limits and permissions.

3. Model Providers & Model Catalog

Located at **Models**  **Model Providers** (</workspace/models/providers>) and **Model Catalog**:

3.1 Connecting Commercial & Private AI Providers

UnifAI supports over 28 providers out of the box:

[illegible]

Step-by-Step Provider Connection:

1. Navigate to **Models**  **Model Providers**.
2. Click **+ Add Provider** and choose the vendor (e.g., OpenAI, Anthropic, Google Gemini, Azure OpenAI, Ollama).
3. Fill in the connection parameters:
 - **Provider Name:** Unique identifier (e.g., **openai-production-tier1**).
 - **API Key:** Third-party vendor API secret (stored with AES-256-GCM encryption).
 - **Base URL (Optional):** Custom proxy URL or private Azure/Ollama endpoint (e.g. **http://ollama-server:11434**).
 - **Weight / Priority:** Integer weight used for automatic weighted load balancing.
4. Click **Test Connection**. The gateway will execute a lightweight health probe to verify credentials.
5. Click **Save Provider**. Available models will automatically synchronize into the **Model Catalog**.

3.2 Fallback Cascades & Adaptive Routing

Configure fallback chains to guarantee 99.999% availability:

- If **gpt-4o** returns an upstream HTTP 500/503 or hits a rate limit, the gateway automatically cascades the request to **claude-3-5-sonnet** or an internal Ollama model with zero client code changes.

4. Virtual Keys & Spend Governance

Located at **Virtual Keys** (/workspace/virtual-keys):

Virtual Keys are the gateway's primary credential mechanism for client applications, developers, and internal systems. They completely shield upstream provider API keys from exposure.

4.1 Creating a Virtual Key

1. Click + **Create Virtual Key**.
2. Provide key parameters:
 - **Key Name:** Identifier (e.g., **customer-support-agent-prod**).

- **Workspace / Team:** Assign to a department or project.
 - **Allowed Models:** Restrict the key to specific models (e.g., only **gpt-4o-mini** and **claude-3-haiku**).
 - **Budget Limit ():** Monthly or total hard dollar spend cap (e.g., **500.00**).
 - **Rate Limits:**
 - Requests Per Minute (RPM) cap (e.g., **120 RPM**).
 - Tokens Per Minute (TPM) cap (e.g., **100,000 TPM**).
3. Click **Generate Key**.
 4. **Copy the generated key immediately.** The secret (**vk-live-...**) is displayed only once.

5. Browser AI / Guard Fleet Management

Located at **Browser AI** (**/workspace/browser-ai/***):

The Browser AI suite allows administrators to manage, monitor, and enforce corporate security policies across all employee web interactions with public AI platforms.

Browser AI Admin Menu ■■■ Target Websites → Configure which AI sites are inspected ■■■ Guard Rules → Define DLP, regex, keywords, and AI Guard Bot rules ■■■ Prompt Logs → Full audit trail of prompts, blocks, and redactions ■■■ Guard Agents → Fleet health, OS versions, and agent status ■■■ Security Insights → Visual analytics on top violations and risky prompts ■■■ Setup & Download → Build, package, and distribute Windows/macOS installers

5.1 Configuring Target Websites

Navigate to **Browser AI** \rightrightarrow **Target Websites**:

1. Click **+ Add Target Website**.
2. Enter the domain or regex pattern:
 - Example 1: **chatgpt.com** (Protocol: HTTPS, Target Name: "OpenAI ChatGPT")
 - Example 2: **claude.ai** (Target Name: "Anthropic Claude")
 - Example 3: **perplexity.ai** (Target Name: "Perplexity AI")
3. Toggle **Enabled** to **Active**.
4. When saved, all connected endpoint guards immediately receive the updated PAC and proxy routing list on their next heartbeat.

5.2 Creating Guard Rules (DLP & Guard Bot)

Navigate to **Browser AI** \rightrightarrow **Guard Rules**:

1. Click **+ Add Rule**.
2. Define Rule Criteria:
 - **Rule Name:** e.g., **Block Credit Card Numbers (PCI-DSS)**.
 - **Target Scope:** Select All Websites or specific target platforms.
 - **Matching Engine:**
 - **Regex:** e.g. **\b(?:4[0-9]{12}(?:[0-9]{3})?|5[1-5][0-9]{14}|3[47][0-9]{13})\b**
 - **Keywords:** Comma-separated list (e.g. **CONFIDENTIAL_PROJECT_APOLLO, merger_terms**).
 - **AI Guard Bot:** Semantic assessment prompt evaluated by local Ollama LLM.

- **Action:**
 - **BLOCK:** Drop the prompt and display an in-browser compliance violation notification.
 - **REDACT:** Mask the sensitive substring with **[REDACTED_PCI_DATA]** and let the prompt proceed.
 - **WARN:** Allow the prompt but flag an alert in security logs and trigger an admin webhook.
3. Set **Priority** (1 – 100, where 1 is highest priority).
 4. Click **Save Rule**.

5.3 Guard Agent Fleet Monitoring

Navigate to **Browser AI**  **Guard Agents**:

- View the active fleet of connected laptops:
- Agent ID, Employee Username, Hostname, Operating System (Windows 11, macOS Sequoia).
- Guard Version (e.g. **v1.1.15**).
- Heartbeat status: **Online** (green indicator, active within 60s) or **Offline**.
- Last synced policy revision.

5.4 Rebuilding & Packaging Guard Installers

Navigate to **Browser AI**  **Setup**:

1. View the current distribution packages for **Windows (EXE)** and **macOS (PKG/ZIP)**.
2. If you updated the server domain or fleet secret, click **Rebuild & Publish Packages**.
3. Download the standalone installers directly or copy the one-line deployment script to distribute via MDM (Microsoft Intune, Jamf Pro, Kandji).
4. **Company Uninstall Key:** Set the mandatory master secret required if an employee or technician attempts to uninstall the guard client.

6. Workspace Configuration Views

UnifAI provides seven dedicated configuration views located under **Workspace**  **Config** (**/workspace/config/***):

6.1 Client Settings (**/workspace/config/client-settings**)

- Configures default timeouts, client retry counts, backoff multipliers, and streaming chunk buffer sizes for all outbound gateway connections.

6.2 Compatibility (**/workspace/config/compatibility**)

- Manages legacy OpenAI API version compatibility layers, Anthropic message formatting adapters, and custom vendor header rewriting rules.

6.3 Caching (**/workspace/config/caching**)

- Configures the In-Memory, Redis, or Vector-based Semantic Cache.
- Set cache TTL (Time to Live in seconds), maximum memory allocation, and similarity thresholds (e.g. **0.92**) for semantic hit matching.

6.4 Security (/workspace/config/security)

- Enforces session expiration timers (default: 24 hours), MFA policies, brute-force lockout thresholds (maximum failed attempts and lockout window in minutes), and trusted proxy CIDRs.

6.5 API Keys (/workspace/config/api-keys)

- Enterprise gate for managing root master API access, key rotation schedules, and third-party credential vaults.

6.6 Performance Tuning (/workspace/config/performance-tuning)

- Fine-tunes fasthttp worker pool size, maximum concurrent connections, TCP keep-alive intervals, and read/write buffer allocations.

6.7 Feature Flags (/workspace/config/feature-flags)

- Toggles runtime system capabilities dynamically without restarting containers:
- **enable_mcp_gateway**: Activates Model Context Protocol proxying.
- **enable_semantic_cache**: Activates vector-based caching engine.
- **enable_ai_guard_bot**: Enables Ollama-powered semantic guardrails.
- **enforce_strict_dlp**: Switches guard fleet to hard fail-closed enforcement.

7. Audit Logging & Security Forensics

Located at **Observability** → **Audit Logs** (/workspace/observability/audit-logs) and **Browser AI** → **Prompt Logs**:

- Every mutating administrative action (key creation, user approval, rule change, provider update) is permanently recorded with timestamp, actor email, source IP address, and JSON diff of changes.
- Prompt logs record every intercepted employee prompt, match scores, triggered DLP rules, and redaction diffs for compliance auditing (HIPAA, SOC 2, ISO 27001, GDPR).

End of Administrator Guide.